



CS4302-01

Parallel and Distributed Computing

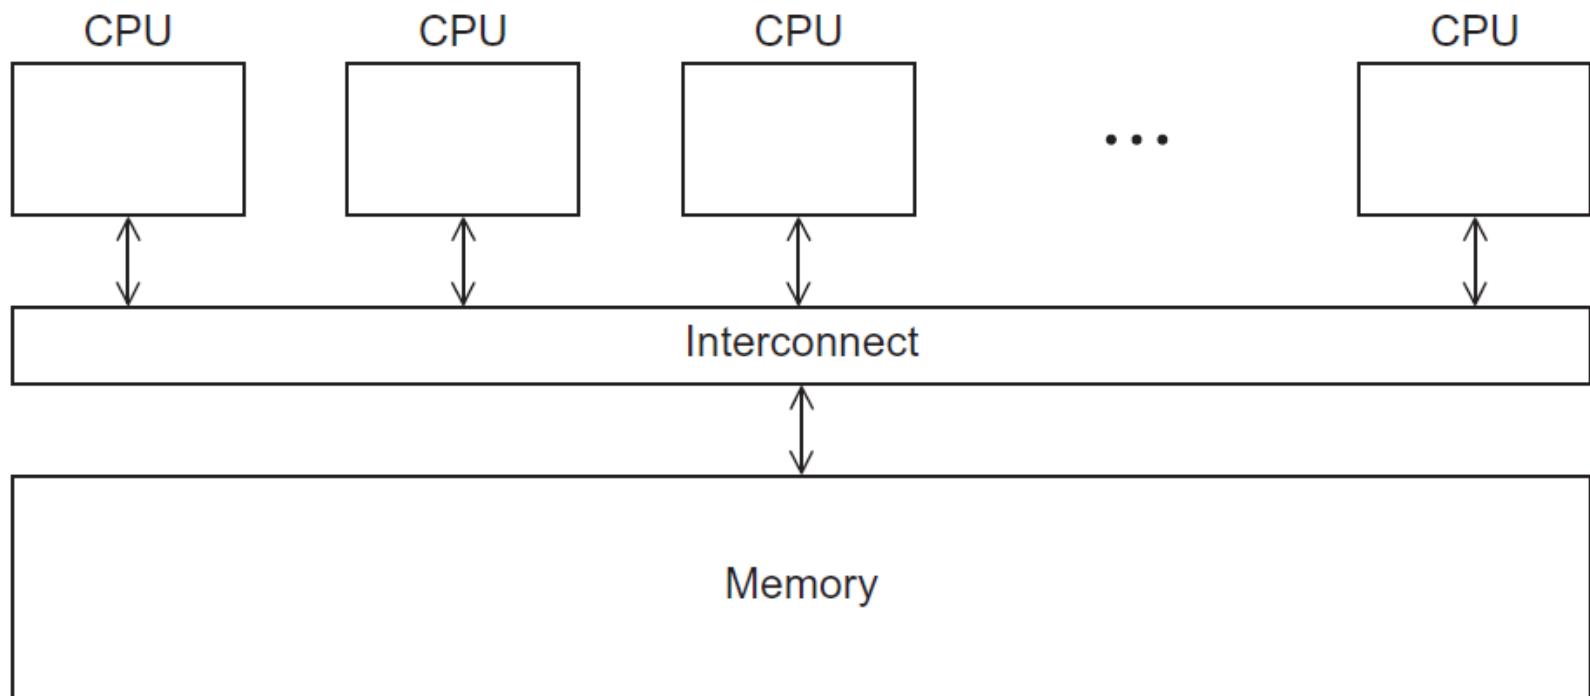
Lecture 3 OpenMP

Zhuoran Song

2023/9/22



A Shared Memory System





Shared Memory Programming

- **Shared Memory Programming**

- *Start a single process and fork threads.*
- *Threads carry out work.*
- *Threads communicate through shared memory.*
- *Threads coordinate through synchronization (also through shared memory).*

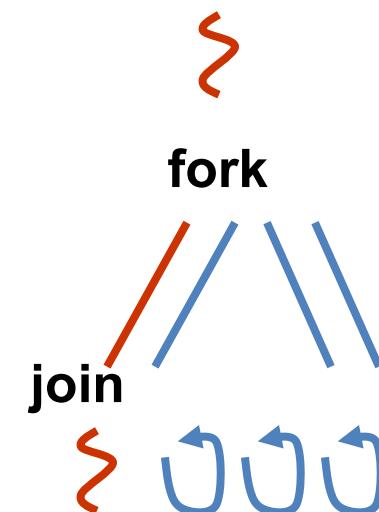
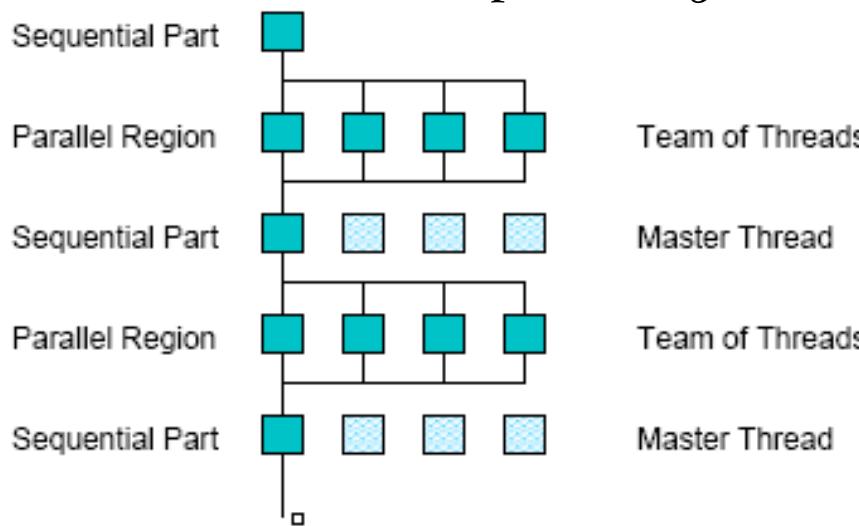
- **Distributed Memory Programming**

- *Start multiple processes on multiple systems.*
- *Processes carry out work.*
- *Processes communicate through message-passing.*
- *Processes coordinate either through message-passing or synchronization (generates messages).*



Execution Model

- Fork-join model of parallel execution
- Begin execution as a single process (**master thread**)
- Start of a parallel construct:
 - Master thread creates team of threads (**worker threads**)
- Completion of a parallel construct:
 - Threads in the team synchronize -- *implicit barrier*
- Only master thread continues execution
- Implementation optimization:
 - Worker threads spin waiting on next fork





OpenMP

- An API for shared-memory parallel programming, MP = multiprocessing
- Designed for systems in which each thread or process can potentially have access to all available memory.
- System is viewed as a collection of cores or CPU's, all of which have access to main memory.
- Higher-level support for scientific programming on shared memory architectures.
- Programmer identifies parallelism and data properties, and guides scheduling at a high level.
- System decomposes parallelism and manages schedule.

See <http://www.openmp.org>



OpenMP

- **Common model for shared-memory parallel programming**
 - *Portable across shared-memory architectures*
- **Scalable (on shared-memory platforms)**
- **Incremental parallelization**
 - *Parallelize individual computations in a program while leaving the rest of the program sequential*
- **Compiler based**
 - *Compiler generates thread program and synchronization*
- **Extensions to existing programming languages (Fortran, C and C++)**
 - *mainly by directives*
 - *a few library routines*



Programmer's View

- **OpenMP is a portable, threaded, shared-memory programming *specification* with “light” syntax**
 - *Exact behavior depends on OpenMP implementation!*
 - *Requires compiler support (C/C++ or Fortran)*
- **OpenMP will:**
 - *Allow a programmer to separate a program into serial regions and parallel regions, rather than concurrently-executing threads.*
 - *Hide stack management*
 - *Provide synchronization constructs*
- **OpenMP will not:**
 - *Parallelize automatically*
 - *Guarantee speedup*
 - *Provide freedom from data races*



Pragmas

- **Pragmas are special preprocessor instructions.**
- **Typically added to a system to allow behaviors that aren't part of the basic C specification.**
- **Compilers that don't support the pragmas ignore them.**
- **The interpretation of OpenMP pragmas**
 - *They modify the statement immediately following the pragma*
 - *This could be a compound statement such as a loop*

#pragma omp ...



“Hello World”

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);    指定多少个线程

# pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);

} /* Hello */
```



“Hello World”

```
gcc -g -Wall -fopenmp -o omp_hello omp_hello.c
```

```
./omp_hello 4
```

compiling

running with 4 threads

Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

possible outcomes

Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4

Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4



Generate Parallel Threads

- **# pragma omp parallel num_threads (thread_count)**
 - Most basic parallel directive.
 - The number of threads that run the following structured block of code is determined by the run-time system.
- **# pragma omp parallel num_threads (thread_count)**
 - Clause: text that modifies a directive.
 - The num_threads clause can be added to a parallel directive.
 - It allows the programmer to specify the number of threads that should execute the following block.



Query Functions

int omp_get_num_threads(void); 目前有多少线程

Returns the number of threads currently in the team executing the parallel region from which it is called

int omp_get_thread_num(void); 目前的线程id

Returns the thread number, within the team, that lies between 0 and `omp_get_num_threads () -1`, inclusive. The master thread of the team is thread 0

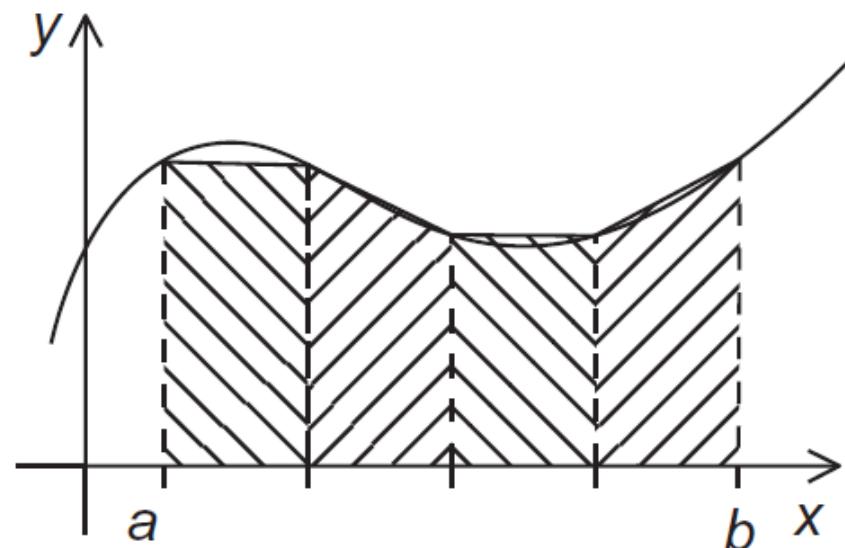
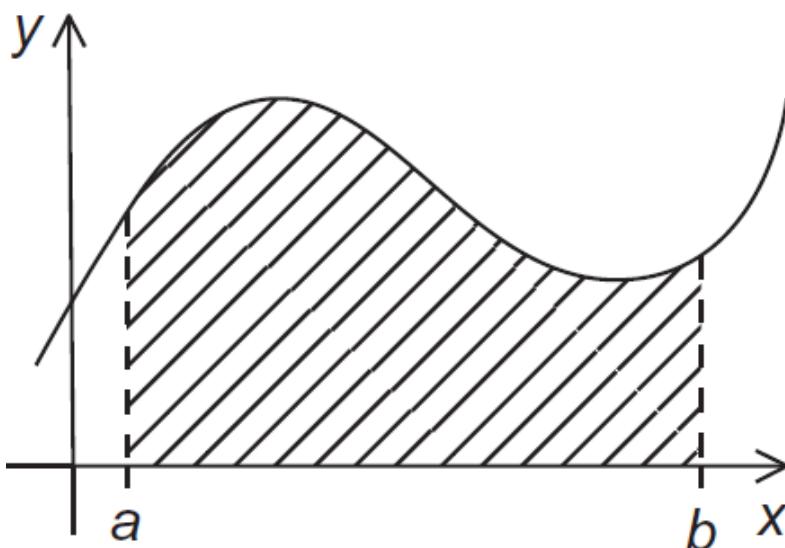
- There may be system-defined limitations on the number of threads that a program can start.

`setenv OMP_NUM_THREADS 16 [csh, tcsh]` 线程上限

- The OpenMP standard doesn't guarantee that this will actually start `thread_count` threads.
- Use the above functions to get the actual thread number and ID.



The Trapezoidal Rule



```
/* Input: a, b, n */  
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++) {  
    x_i = a + i*h;  
    approx += f(x_i);  
}  
approx = h*approx;
```

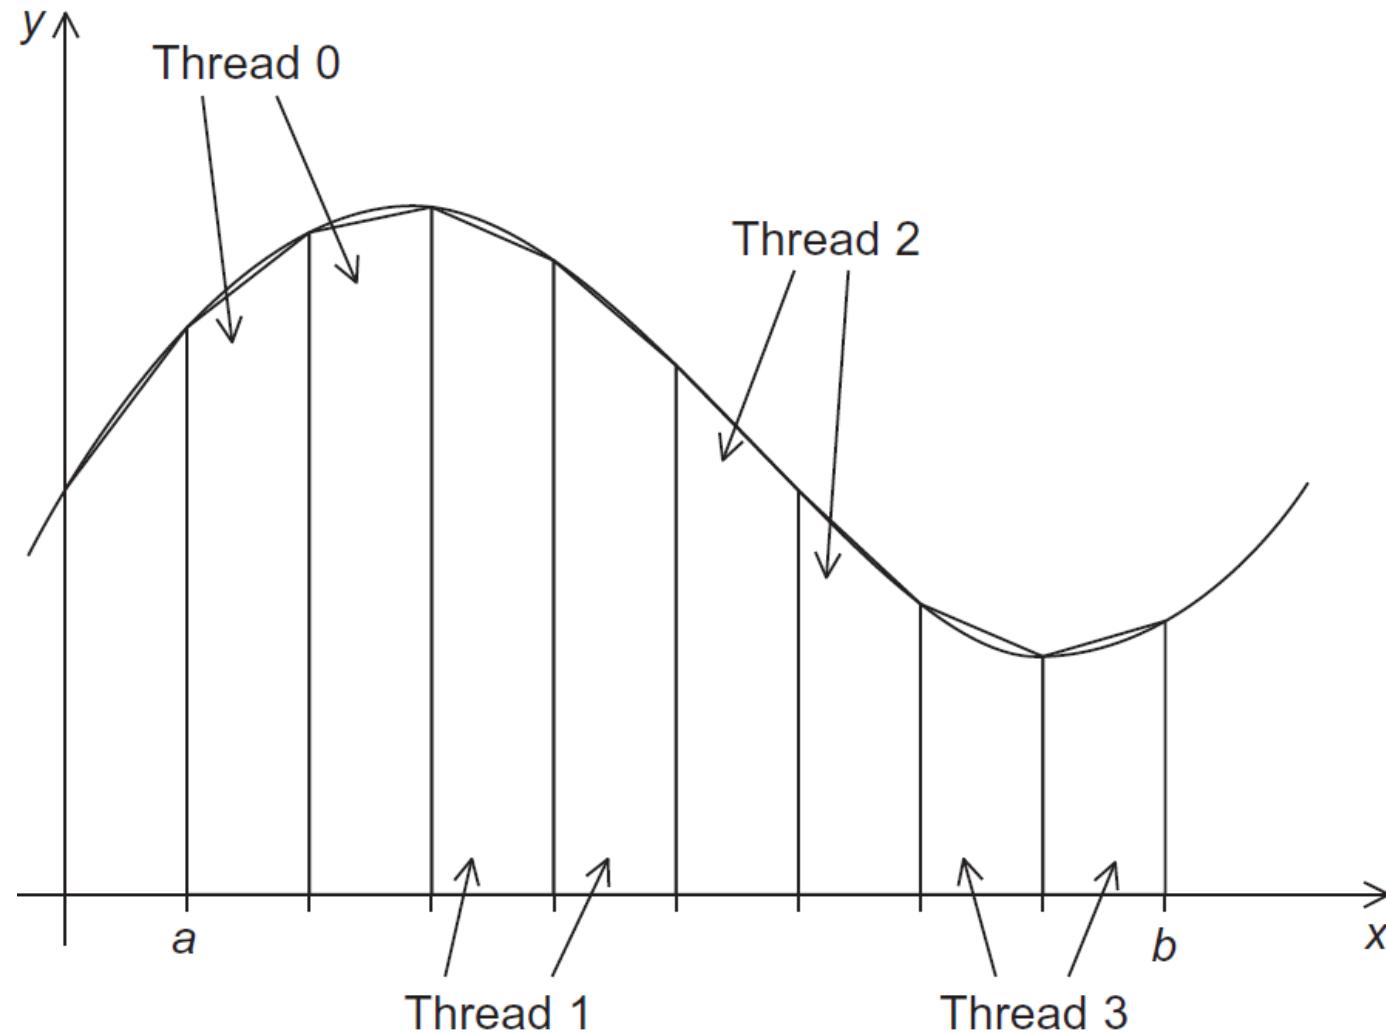


OpenMP Implementation

- **We identified two types of tasks:**
 - a) computation of the areas of individual trapezoids, and*
 - b) adding the areas of trapezoids.*
- **There is no communication among the tasks in the first collection, but each task in the first collection communicates with task b.**
- **So we aggregated tasks by assigning a contiguous block of trapezoids to each thread (and a single thread to each core).**



OpenMP Implementation





OpenMP Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Trap(double a, double b, int n, double* global_result_p);

int main(int argc, char* argv[]) {
    double global_result = 0.0; /* Store result in global_result */
    double a, b;               /* Left and right endpoints */
    int n;                     /* Total number of trapezoids */
    int thread_count;

    thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);
# pragma omp parallel num_threads(thread_count)
    Trap(a, b, n, &global_result);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n",
           a, b, global_result);
    return 0;
} /* main */
```



OpenMP Implementation

```
void Trap(double a, double b, int n, double* global_result_p) {
    double h, x, my_result;
    double local_a, local_b;
    int i, local_n;
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    h = (b-a)/n;
    local_n = n/thread_count;
    local_a = a + my_rank*local_n*h;
    local_b = local_a + local_n*h;
    my_result = (f(local_a) + f(local_b))/2.0;
    for (i = 1; i <= local_n-1; i++) {
        x = local_a + i*h;
        my_result += f(x);
    }
    my_result = my_result*h;

    # pragma omp critical
    *global_result_p += my_result;
}. /* Trap */
```



Mutual Exclusion

Time	Thread 0	Thread 1
0	global_result = 0 to register	finish my_result
1	my_result = 1 to register	global_result = 0 to register
2	add my_result to global_result	my_result = 2 to register
3	store global_result = 1	add my_result to global_result
4		store global_result = 2

Unpredictable results when two (or more) threads attempt to simultaneously execute:

global_result += my_result ;



Reduction

- A **reduction operator** is a binary operation (such as addition or multiplication).
- A **reduction** is a computation that repeatedly applies the same reduction operator to a sequence of operands in order to get a single result.
- All of the intermediate results of the operation should be stored in the same variable: the reduction variable.

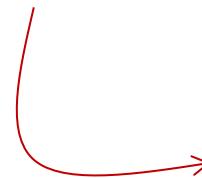
+ **bitwise & logical &**
- **bitwise | logical |**
* **bitwise ^ max/min**



Reduction

A reduction clause can be added to a parallel directive.

```
reduction(<operator>: <variable list>)
```



+, *, -, &, |, ^, &&, ||

```
global_result = 0.0;  
# pragma omp parallel num_threads(thread_count) \  
    reduction(+: global_result)  
    global_result += Local_trap(double a, double b, int n);
```



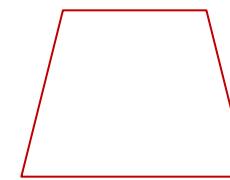
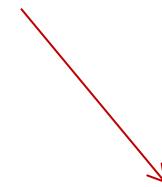
“Parallel For”

- Forks a team of threads to execute the following structured block.
- However, the structured block following the parallel for directive must be a for loop.
- Furthermore, with the parallel for directive the system parallelizes the for loop by dividing the iterations of the loop among the threads.



“Parallel For”

```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```



```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+: approx)  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```



Caveats

```
fibo[ 0 ] = fibo[ 1 ] = 1;  
for (i = 2; i < n; i++)  
    fibo[ i ] = fibo[ i - 1 ] + fibo[ i - 2 ];
```



```
fibo[ 0 ] = fibo[ 1 ] = 1;
```

note 2 threads

```
# pragma omp parallel for num_threads(2)
```

```
for (i = 2; i < n; i++)
```

```
    fibo[ i ] = fibo[ i - 1 ] + fibo[ i - 2 ];
```



1 1 2 3 5 8 13 21 34 55

this is correct

but sometimes
we get this

1 1 2 3 5 8 0 0 0 0





Caveats

- OpenMP compilers don't check for dependences among iterations in a loop that's being parallelized with a parallel for directive.
- A loop in which the results of one or more iterations depend on other iterations cannot, in general, be correctly parallelized by OpenMP.



Estimating PI

$$\pi = 4 \left[1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

```
double factor = 1.0;
double sum = 0.0;
for (k = 0; k < n; k++) {
    sum += factor/(2*k+1);
    factor = -factor;
}
pi_approx = 4.0*sum;
```



Estimating PI

loop dependency

```
double factor = 1.0;
double sum = 0.0;
#pragma omp parallel for num_threads(thread_count) \
    reduction(+:sum)
for (k = 0; k < n; k++) {
    sum += factor/(2*k+1);
    factor = -factor;
}
pi_approx = 4.0*sum;
```



Scope of Variables

- A variable that can be accessed by all the threads in the team has **shared** scope.
- A variable that can only be accessed by a single thread has **private** scope.
- The default scope for variables declared before a parallel block is **shared**.



Estimating PI

```
double sum = 0.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+:sum) private(factor)  
for (k = 0; k < n; k++) {  
    if (k % 2 == 0)  
        factor = 1.0;  
    else  
        factor = -1.0;  
    sum += factor/(2*k+1);  
}
```



Insures factor has
private scope.



Scope of Variables

- Lets the programmer specify the scope of each variable in a block. **default**(none)
- With this clause the compiler will require that we specify the scope of each variable we use in the block and that has been declared outside the block.

```
#include <omp.h>
double sum = 0.0;
#pragma omp parallel for num_threads(thread_count) \
    default(none) reduction(+:sum) private(k, factor) \
    shared(n)
for (k = 0; k < n; k++) {
    if (k % 2 == 0)
        factor = 1.0;
    else
        factor = -1.0;
    sum += factor/(2*k+1);
}
```



Loop Carried Dependence

```
for (i=0; i<100; i++) {  
    A[i+1]=A[i]+C[i];  
    B[i+1]=B[i]+A[i+1];  
}
```

- **Two loop carried dependence.**
- **One intra-loop dependence.**



Loop Carried Dependence

```
for (i=0; i<100; i++) {  
    A[i]=A[i]+B[i];  
    B[i+1]=C[i]+D[i];  
}
```

Eliminating loop dependence:

```
A[0]=A[0]+B[0];  
for (i=0; i<99; i++) {  
    B[i+1]=C[i]+D[i];  
    A[i+1]=A[i+1]+B[i+1];  
}  
B[100]=C[99]+D[99];
```



Loop Scheduling

- Default schedule:

```
sum = 0.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+:sum)  
for (i = 0; i <= n; i++)  
    sum += f(i);
```

- Cyclic schedule:

```
sum = 0.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+:sum) schedule(static,1)  
for (i = 0; i <= n; i++)  
    sum += f(i);
```



Loop Scheduling

schedule (type , chunkszie)

- Type can be:
 - *static*: *the iterations can be assigned to the threads before the loop is executed.*
 - *dynamic or guided*: *the iterations are assigned to the threads while the loop is executing.*
 - *auto*: *the compiler and/or the run-time system determine the schedule.*
 - *runtime*: *the schedule is determined at run-time.*
- The chunkszie is a positive integer.



Static Scheduling

schedule(static, 1)

.....

Thread 0 : 0,3,6,9

Thread 1 : 1,4,7,10

Thread 2 : 2,5,8,11

schedule(static, 2)

Thread 0 : 0,1,6,7

Thread 1 : 2,3,8,9

Thread 2 : 4,5,10,11

schedule(static, 4)

Thread 0 : 0,1,2,3

Thread 1 : 4,5,6,7

Thread 2 : 8,9,10,11



Dynamic Scheduling

- The iterations are also broken up into chunks of **chunksize** consecutive iterations.
- Each thread executes a chunk, and when a thread finishes a chunk, it requests another one from the run-time system.
- This continues until all the iterations are completed.
- The **chunksize** can be omitted. When it is omitted, a **chunksize** of 1 is used.



Guided Scheduling

- Each thread also executes a chunk, and when a thread finishes a chunk, it requests another one.
- However, in a guided schedule, as chunks are completed the size of the new chunks decreases exponentially.
- If no **chunksize** is specified, the size of the chunks decreases down to 1.
- If **chunksize** is specified, it decreases down to **chunksize**, with the exception that the very last chunk can be smaller than **chunksize**.



Guided Scheduling

Thread	Chunk	Size of Chunk	Remaining Iterations
0	1 – 5000	5000	4999
1	5001 – 7500	2500	2499
1	7501 – 8750	1250	1249
1	8751 – 9375	625	624
0	9376 – 9687	312	312
1	9688 – 9843	156	156
0	9844 – 9921	78	78
1	9922 – 9960	39	39
1	9961 – 9980	20	19
1	9981 – 9990	10	9
1	9991 – 9995	5	4
0	9996 – 9997	2	2
1	9998 – 9998	1	1
0	9999 – 9999	1	0

**Assignment of trapezoidal rule iterations 1–9999
using a guided schedule with two threads.**



Runtime Scheduling

- The system uses the environment variable **OMP_SCHEDULE** to determine at run-time how to schedule the loop.
- The **OMP_SCHEDULE** environment variable can take on any of the values that can be used for a static, dynamic, or guided schedule.
- For example
 - setenv OMP_SCHEDULE GUIDED,4 [csh, tcsh]
 - export OMP_SCHEDULE=GUIDED,4 [sh, ksh, bash]



Impact of Scheduling

- Load balance
 - *Same work in each iteration?*
 - *Processors working at same speed?*
- Scheduling overhead
 - *Static decisions are cheap because they require no run-time coordination*
 - *Dynamic decisions have overhead that is impacted by complexity and frequency of decisions*
- Data locality
 - *Particularly within cache lines for small chunk sizes*
 - *Also impacts data reuse on same processor*



Data Locality

- Consider a 1-Dimensional array to solve the global sum problem, 16 elements, 4 threads

CYCLIC (chunk = 1):



BLOCK (chunk = 4):





Data Locality

Consider how data is accessed

- **Temporal locality**
 - *Same or nearby data used multiple times*
 - *Intrinsic in computation*
- **Spacial locality**
 - *Data nearby to be used and is present in “fast memory”*
 - *Same data transfer (same cache line, same DRAM transaction)*
- **What can we do to get locality**
 - *Appropriate data placement and layout*
 - *Code reordering transformations*